

TensorFlow Tutorial #02

Classification

Seoung Wug Oh, Ph.D. Student
Computational Intelligence and Photography Lab.
Yonsei University

Classification

- Binary classification (logistic regression)
 - **Email:** Spam or Not ?
 - **Tumor:** Malignant or Benign ?

$$y \in \{0, 1\}$$


- Multi-class classification (softmax regression)
 - **Image:** dog ? Cup ? Hat ?
 - **Human Activity:** running ? Walking ? Clapping ?
 - **Digits:** 1? 2? 3? ...

$$y \in \{0, 1, 2, 3, \dots, K\}$$

Logistic regression


- Logistic Regression
 - Given input x , model that can estimate y .
 - Want $\hat{y} = f(x)$, where $0 \leq \hat{y} \leq 1$.
 - Function f can be in any form.
 - In this session, f is simple linear function.
- Linear Model
 - Relate input x and output \hat{y} **linearly**
$$\text{logit} = \mathbf{W} \times x + b$$
 - *logit* is unbounded value.
 - Higher values mean higher probability of being positive (1), and vice versa.

Logistic regression




Logistic regression

- Sigmoid function
 - Squeeze *logit* to output $0 \leq \hat{y} \leq 1$.



Logistic regression


$$\text{sig}(t) = \frac{1}{1+e^{-t}}$$

- Final model
 - $\hat{y} = f(x) = \text{sig}(W \times x + b)$
 - Goal:
Optimize parameter set $\theta = \{W, b\}$ to fit $\hat{y} \approx y$ where y is ground-truth label.
 - If $f(x)$ is spam detector, $f(x) = 0.7$ tells that 70% chance of the mail is being spam mail.
- Binary cross entropy
 - $C = y \times -\log(\hat{y}) + (1 - y) \times -\log(1 - \hat{y})$,
where $y \in \{0, 1\}$

Example: Tumor classification

- Problem
 - Given some features, estimate the tumor is whether malignant or benign
- Dataset: Wisconsin Diagnostic Breast Cancer (WDBC)
 - $x = \text{features of tumor}$ (radius, texture, etc)
 - $y \in \{\text{malignant}, \text{benign}\}$

id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness
842302	M	17.99	10.38	122.8	1001	0.1184
842517	M	20.57	17.77	132.9	1326	0.08474
84300903	M	19.69	21.25	130	1203	0.1096
84348301	M	11.42	20.38	77.58	386.1	0.1425

Build Computational Graph

```
# Place holders  
x = tf.placeholder(tf.float32, [None, 30]) # Inputs: a batch of features (30 dims)  
y = tf.placeholder(tf.float32, [None, 1]) # Labels: a batch of labels
```


Build Computational Graph

Place holders

```
x = tf.placeholder(tf.float32, [None, 30]) # Inputs: a batch of features (30 dims)
y = tf.placeholder(tf.float32, [None, 1]) # Labels: a batch of labels
```

Set model weights

```
W = tf.Variable(tf.zeros([30, 1]))
b = tf.Variable(tf.zeros([1]))
```

Build Computational Graph

```
# Place holders
```

```
x = tf.placeholder(tf.float32, [None, 30]) # Inputs: a batch of features (30 dims)  
y = tf.placeholder(tf.float32, [None, 1]) # Labels: a batch of labels
```

```
# Set model weights
```

```
W = tf.Variable(tf.zeros([30, 1]))  
b = tf.Variable(tf.zeros([1]))
```

```
# Construct model ( $y = W \cdot X + b$ )
```

```
logit = tf.matmul(x, W) + b  
pred = tf.nn.sigmoid(logit)
```

Build Computational Graph

Place holders

```
x = tf.placeholder(tf.float32, [None, 30]) # Inputs: a batch of features (30 dims)
y = tf.placeholder(tf.float32, [None, 1]) # Labels: a batch of labels
```

Set model weights

```
W = tf.Variable(tf.zeros([30, 1]))
b = tf.Variable(tf.zeros([1]))
```

Construct model ($y = W \cdot X + b$)

```
logit = tf.matmul(x, W) + b
pred = tf.nn.sigmoid(logit)
```

Directly compute loss from logit (to ensure stability and avoid overflow)

```
cost = tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(logits=logit,
labels=y))
```

Build Computational Graph

Place holders

```
x = tf.placeholder(tf.float32, [None, 30]) # Inputs: a batch of features (30 dims)
y = tf.placeholder(tf.float32, [None, 1]) # Labels: a batch of labels
```

Set model weights

```
W = tf.Variable(tf.zeros([30, 1]))
b = tf.Variable(tf.zeros([1]))
```

Construct model ($y = W \cdot X + b$)

```
logit = tf.matmul(x, W) + b
pred = tf.nn.sigmoid(logit)
```

Directly compute loss from logit (to ensure stability and avoid overflow)

```
cost = tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(logits=logit,
labels=y))
```

Define optimizer and train_op

```
train_op = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

Implementation Note

- Points
 - `tf.nn.sigmoid_cross_entropy_with_logits(logits=logit, labels=y)`

```
z * -log(sigmoid(x)) + (1 - z) * -log(1 - sigmoid(x))  
= z * -log(1 / (1 + exp(-x))) + (1 - z) * -log(exp(-x) / (1 + exp(-x)))  
= z * log(1 + exp(-x)) + (1 - z) * (-log(exp(-x)) + log(1 + exp(-x)))  
= z * log(1 + exp(-x)) + (1 - z) * (x + log(1 + exp(-x)))  
= (1 - z) * x + log(1 + exp(-x))  
= x - x * z + log(1 + exp(-x))
```

For $x < 0$, to avoid overflow in $\exp(-x)$, we reformulate the above

```
x - x * z + log(1 + exp(-x))  
= log(exp(x)) - x * z + log(1 + exp(-x))  
= -x * z + log(1 + exp(x))
```

Hence, to ensure stability and avoid overflow, the implementation uses this equivalent formulation

```
max(x, 0) - x * z + log(1 + exp(-abs(x)))
```

Training Loop

```
with tf.Session() as sess:  
    # Initializing the variables  
    sess.run(tf.global_variables_initializer())
```

Training Loop

```
with tf.Session() as sess:
    # Initializing the variables
    sess.run(tf.global_variables_initializer())

    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0.
        total_batch = int(X_train.shape[0]/batch_size)
        # Loop over all batches
        for i in range(total_batch):
            batch_xs = X_train[i:i+batch_size,:]
            batch_ys = y_train[i:i+batch_size,:]
            # Run optimization op (backprop) and cost op (to get loss value)
            _, c = sess.run([train_op, cost], feed_dict={x: batch_xs, y: batch_ys})
            # Compute average loss
            avg_cost += c / total_batch
```

Training Loop

```
with tf.Session() as sess:
    # Initializing the variables
    sess.run(tf.global_variables_initializer())

    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0.
        total_batch = int(X_train.shape[0]/batch_size)
        # Loop over all batches
        for i in range(total_batch):
            batch_xs = X_train[i:i+batch_size,:]
            batch_ys = y_train[i:i+batch_size,:]
            # Run optimization op (backprop) and cost op (to get loss value)
            _, c = sess.run([train_op, cost], feed_dict={x: batch_xs, y: batch_ys})
            # Compute average loss
            avg_cost += c / total_batch

        # Display logs per epoch step
        if (epoch+1) % display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "cost=", "{:.9f}".format(avg_cost))
```


Testing

```
# Test model
```

```
correct_prediction = tf.equal(tf.round(pred), y)
```

```
# Calculate accuracy
```

```
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```

```
print("Accuracy:", accuracy.eval({x: X_test, y: y_test}))
```

Testing

```
# Test model
correct_prediction = tf.equal(tf.round(pred), y)

# Calculate accuracy
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print("Accuracy:", accuracy.eval({x: X_test, y: y_test}))
```

- Classification
 - Decision boundary = 0.5 (rounding)
 - Accuracy the percentage of correct estimation

```
accuracy.eval({x: X_test, y: y_test}))
sess.run(accuracy, feed_dict={x: X_test, y: y_test})
```

Results

```
> > python logistic_regression.py
```

```
...
```

```
Epoch: 0095 cost= 0.024283896
```

```
Epoch: 0096 cost= 0.024087359
```

```
Epoch: 0097 cost= 0.023894221
```

```
Epoch: 0098 cost= 0.023704397
```

```
Epoch: 0099 cost= 0.023517799
```

```
Epoch: 0100 cost= 0.023334336
```

```
Optimization Finished!
```

```
Accuracy: 0.913043
```

Multi-class Classification (softmax regression)

- Binary classification (logistic regression)
 - Email: Spam or Not ?
 - Tumor: Malignant or Benign ?


$$y \in \{0, 1\}$$

- **Multi-class classification (softmax regression)**
 - Image: dog ? Cup ? Hat ?
 - Human Activity: running ? Walking ? Clapping ?
 - Digits: 1? 2? 3? ...

$$y \in \{0, 1, 2, 3, \dots, K\}$$


Softmax regression

- Logistic Regression
 - $\hat{y} = f(x)$, where $0 \leq \hat{y} \leq 1$.
 - Negative or Positive
- Softmax regression
 - Now, $\hat{y} = f(x)$ is the **probability distribution** over classes, where $0 \leq \hat{y} \leq 1$ and $\sum \hat{y} = 1$.
 - $\text{logit} \in \mathbb{R}^K = \mathbf{W} \times x + b$
 - Higher values at n -th element mean higher probability of being belong to n -th class



Softmax regression



- Logistic Regression



Softmax regression

- Softmax function

$$\sigma(x_j) = \frac{e^{x_j}}{\sum_i e^{x_i}}$$




Softmax regression

$$\sigma(x_j) = \frac{e^{x_j}}{\sum_i e^{x_i}}$$

- Final model
 - $\hat{y} = f(x) = \textit{softmax}(W \times x + b)$
 - Goal:
Optimize parameter set $\theta = \{W, b\}$ to fit $\hat{y} \approx y$ where y is ground-truth label.
- Multinomial cross entropy
 - $C = y^k \times -\log(\hat{y}^k)$
where y is one-hot vector that indicate the ground-truth class
 - E.g.) $y \in \{0,0,0,0,1,0,0,0,0,0\}$

Example: 2D classification

- Problem
 - Given 2D data, learn a classifier
- Dataset: randomly generated labeled 2D data
 - x = 2D coordinate
 - y = label



Build TF Graph

```
# Place holders  
x = tf.placeholder(tf.float32, [None, 2]) # 2 dimensional input  
y = tf.placeholder(tf.float32, [None, 3]) # 3 classes
```

Build TF Graph

```
# Place holders
x = tf.placeholder(tf.float32, [None, 2]) # 2 dimensional input
y = tf.placeholder(tf.float32, [None, 3]) # 3 classes

# Set model weights
W = tf.Variable(tf.zeros([2, 3]))
b = tf.Variable(tf.zeros([3]))
```

Build TF Graph

Place holders

```
x = tf.placeholder(tf.float32, [None, 2]) # 2 dimensional input  
y = tf.placeholder(tf.float32, [None, 3]) # 3 classes
```

Set model weights

```
W = tf.Variable(tf.zeros([2, 3]))  
b = tf.Variable(tf.zeros([3]))
```

Construct model

```
logit = tf.matmul(x, W) + b  
pred = tf.nn.softmax(logit) # Softmax
```

Build TF Graph

Place holders

```
x = tf.placeholder(tf.float32, [None, 2]) # 2 dimensional input  
y = tf.placeholder(tf.float32, [None, 3]) # 3 classes
```

Set model weights

```
W = tf.Variable(tf.zeros([2, 3]))  
b = tf.Variable(tf.zeros([3]))
```

Construct model

```
logit = tf.matmul(x, W) + b  
pred = tf.nn.softmax(logit) # Softmax
```

Directly compute loss from logit (to ensure stability and avoid overflow)

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logit,  
labels=y))
```

Build TF Graph

Place holders

```
x = tf.placeholder(tf.float32, [None, 2]) # 2 dimensional input  
y = tf.placeholder(tf.float32, [None, 3]) # 3 classes
```

Set model weights

```
W = tf.Variable(tf.zeros([2, 3]))  
b = tf.Variable(tf.zeros([3]))
```

Construct model

```
logit = tf.matmul(x, W) + b  
pred = tf.nn.softmax(logit) # Softmax
```

Directly compute loss from logit (to ensure stability and avoid overflow)

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logit,  
labels=y))
```

Define optimizer and train_op

```
train_op = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

Training Loop (same as logistic regression)

```
with tf.Session() as sess:
    # Initializing the variables
    sess.run(tf.global_variables_initializer())

    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0.
        total_batch = int(X_train.shape[0]/batch_size)
        # Loop over all batches
        for i in range(total_batch):
            batch_xs = X_train[i:i+batch_size,:]
            batch_ys = y_train[i:i+batch_size,:]
            # Run optimization op (backprop) and cost op (to get loss value)
            _, c = sess.run([train_op, cost], feed_dict={x: batch_xs, y: batch_ys})
            # Compute average loss
            avg_cost += c / total_batch

        # Display logs per epoch step
        if (epoch+1) % display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "cost=", "{:.9f}".format(avg_cost))
```

Result

- Training log

Epoch: 0244 cost= 0.254794386

Epoch: 0245 cost= 0.254382825

Epoch: 0246 cost= 0.253973836

Epoch: 0247 cost= 0.253567417

Epoch: 0248 cost= 0.253163518

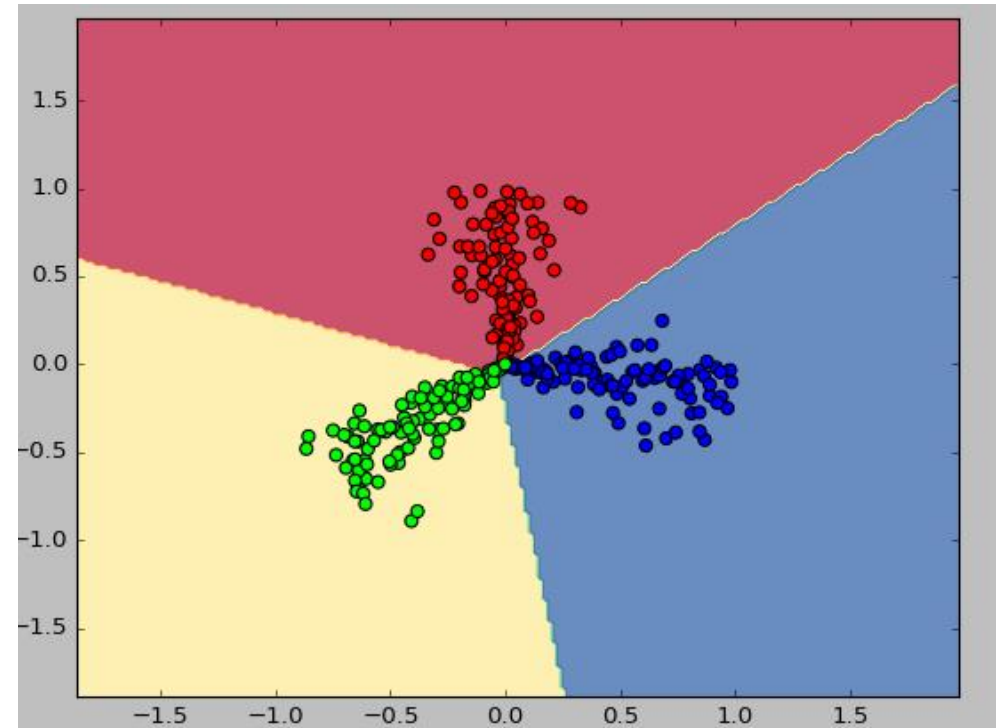


Epoch: 0249 cost= 0.252762134

Epoch: 0250 cost= 0.252363236

Optimization Finished!

Result


- Decision boundary



Practice: 2D spiral data


- Starter Code is provided:
 - Session1_softmax_regression_2D_spiral_starter.py
- Build TF graph your own

```
##### Build graph #####  
# Write your code here  
#####
```




Practice: 2D spiral data

- Limitation of linear model



Practice: 2D spiral data

- Let's move on to Neural Network !



Acknowledgement

- Stanford CS231n
 - <http://cs231n.stanford.edu/>
 - <http://cs231n.github.io/>
- Andrew Ng's ML course
 - <https://www.coursera.org/learn/machine-learning>
- 모두를 위한 머신러닝/딥러닝 강의
 - <https://hunkim.github.io/ml/>